

## Install Package

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import normalize
from sklearn import ensemble
from sklearn.feature_selection import VarianceThreshold
from sklearn.manifold import TSNE
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score
from matplotlib import pyplot
```

```
In [2]: import xgboost as xgb
```

C:\Users\ADMIN\Anaconda3\lib\site-packages\dask\config.py:168: YAMLLoadWarning: calling yaml.load() without Loader=... is deprecated, as the default Loader is unsafe. Please read <https://msg.pyyaml.org/load> for full details.

```
data = yaml.load(f.read()) or {}
```

C:\Users\ADMIN\Anaconda3\lib\site-packages\dask\dataframe\utils.py:13: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.

```
import pandas.util.testing as tm
```

C:\Users\ADMIN\Anaconda3\lib\site-packages\distributed\config.py:20: YAMLLoadWarning: calling yaml.load() without Loader=... is deprecated, as the default Loader is unsafe. Please read <https://msg.pyyaml.org/load> for full details.

```
defaults = yaml.load(f)
```

```
In [3]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score
```

## Load Data

all missing values were removed. Other missing values are imputed using missRanger() in R.

```
In [4]: traindf = pd.read_csv("D:/Github/Risk Modelling/Home_Credit/imputed_application_Train3.csv", index_col=0)
```

```
In [84]: testdf = pd.read_csv("D:/Github/Risk Modelling/Home_Credit/imputed_application_Test3.csv", index_col=0)
```

```
In [6]: traindf.drop('SK_ID_CURR', axis=1, inplace=True)
        #testdf.drop('SK_ID_CURR', axis=1, inplace=True)
```

```
In [7]: traindf.reset_index(drop=True, inplace=True)
```

```
In [63]: testdf.reset_index(drop=True, inplace=True)
```

Check the shape of Train and Test

```
In [9]: traindf.shape
```

```
Out[9]: (307511, 121)
```

```
In [10]: testdf.shape
```

```
Out[10]: (48744, 121)
```

Check If Train and Test still have any Missing Values.

```
In [11]: traindf.isnull().sum().sum()
```

```
Out[11]: 0
```

```
In [12]: testdf.isnull().sum().sum()
```

```
Out[12]: 0
```

```
In [13]: Y_train = traindf['TARGET']
        traindf.drop('TARGET', axis=1, inplace=True)
```

## PreProcessing

### Dimensionality Reduction / Feature Selection

```
In [14]: numerics = ['int16', 'int32', 'int64', 'float16', 'float32', 'float64']
        train_numerical = traindf.select_dtypes(include=numerics)
```

### Correlation


```
In [15]: # Calculate the correlation matrix and take the absolute value
corr_matrix = train_numerical.corr().abs()
# Create a True/False mask and apply it
mask_highcor = np.triu(np.ones_like(corr_matrix, dtype=bool))
tri_df = corr_matrix.mask(mask_highcor)
# List column names of highly correlated features (r > 0.95)
to_drop_corr = [c for c in tri_df.columns if any(tri_df[c] > 0.95)]
to_drop_corr
```

```
Out[15]: ['AMT_CREDIT',
'DAYS_EMPLOYED',
'REGION_RATING_CLIENT',
'APARTMENTS_AVG',
'BASEMENTAREA_AVG',
'YEARS_BEGINEXPLUATATION_AVG',
'YEARS_BUILD_AVG',
'COMMONAREA_AVG',
'ELEVATORS_AVG',
'ENTRANCES_AVG',
'FLOORSMAX_AVG',
'FLOORSMIN_AVG',
'LANDAREA_AVG',
'LIVINGAPARTMENTS_AVG',
'LIVINGAREA_AVG',
'NONLIVINGAPARTMENTS_AVG',
'NONLIVINGAREA_AVG',
'APARTMENTS_MODE',
'BASEMENTAREA_MODE',
'YEARS_BEGINEXPLUATATION_MODE',
'YEARS_BUILD_MODE',
'COMMONAREA_MODE',
'ELEVATORS_MODE',
'ENTRANCES_MODE',
'FLOORSMAX_MODE',
'FLOORSMIN_MODE',
'LANDAREA_MODE',
'LIVINGAPARTMENTS_MODE',
'LIVINGAREA_MODE',
'NONLIVINGAPARTMENTS_MODE',
'NONLIVINGAREA_MODE',
'APARTMENTS_MEDI',
'OBS_30_CNT_SOCIAL_CIRCLE']
```

## Check If Dataset is Imbalance

```
In [16]: Y_train.value_counts().plot.bar()
plt.xlabel('Class')
plt.ylabel('Frequency')
Y_train.value_counts()
```

```
Out[16]: 0    282686
         1    24825
         Name: TARGET, dtype: int64
```



### Random Forest for Feature Selection (aka. Variable Importance)

```
In [17]: # Fit the random forest model to the training data
rf = ensemble.RandomForestClassifier(random_state=123, n_jobs = -1)
rf.fit(train_numerical, Y_train)
```

```
Out[17]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=-1, oob_score=False, random_state=123, verbose=
                                0,
                                warm_start=False)
```

```
In [18]: mask_rf = rf.feature_importances_ > 0.1
          mask_rf
```

[illegible]

Variable Importance could not provide useful info for this Dataset.

## Low Variance Features

```
In [19]: train_numerical_normalized = normalize(train_numerical)
```

```
In [20]: train_numerical_normalized = pd.DataFrame(train_numerical_normalized, columns=
train_numerical.columns)
```

```
In [21]: train_numerical_normalized.describe()
```

Out[21]:

	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY	AMT_GOODS_PRICE
count	3.075110e+05	307511.000000	307511.000000	307511.000000	307511.000000
mean	7.582182e-07	0.250196	0.680537	0.035601	0.609699
std	1.718293e-06	0.154356	0.097865	0.013674	0.089199
min	0.000000e+00	0.008285	0.004808	0.000224	0.003889
25%	0.000000e+00	0.137588	0.662149	0.025049	0.580669
50%	0.000000e+00	0.209651	0.701143	0.032915	0.633069
75%	9.049948e-07	0.321184	0.740990	0.041307	0.665529
max	8.554940e-05	0.999981	0.945914	0.088115	0.984389

8 rows × 104 columns



```
In [26]: Categorical_Level =traindf[categorical_columns].nunique().sort_values(ascending=False)
Categorical_Level
```

```
Out[26]: ORGANIZATION_TYPE          57
OCCUPATION_TYPE          18
NAME_INCOME_TYPE          8
WALLSMATERIAL_MODE        7
WEEKDAY_APPR_PROCESS_START  7
NAME_TYPE_SUITE            7
NAME_HOUSING_TYPE          6
NAME_FAMILY_STATUS          6
NAME_EDUCATION_TYPE         5
FONDKAPREMONT_MODE         4
HOUSETYPE_MODE             3
CODE_GENDER                3
EMERGENCYSTATE_MODE        2
FLAG_OWN_REALTY             2
FLAG_OWN_CAR                2
NAME_CONTRACT_TYPE          2
dtype: int64
```

If the Variable has more than 5 levels then It would be applied LabelEncoder, otherwise applied OHE

```
In [27]: OHE_List = Categorical_Level[Categorical_Level<=5].index.tolist()
LE_List = Categorical_Level[Categorical_Level>5].index.tolist()
```

```
In [28]: le = LabelEncoder()
# Apply LabelEncoder to categorical columns
df_le = traindf[LE_List].apply(lambda x: le.fit_transform(x))
```

```
In [29]: df_ohe = pd.get_dummies(traindf[OHE_List])
```

## Train Data after Converting

```
In [33]: traindf.drop(categorical_columns, axis=1, inplace=True)
```


```
In [34]: traindf = pd.concat([traindf,df_ohe, df_le], axis=1)
```

In [35]: `traindf.head()`

Out[35]:

	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY	AMT_GOODS_PRICE	RE
0	0	202500.0	406597.5	24700.5	351000.0	
1	0	270000.0	1293502.5	35698.5	1129500.0	
2	0	67500.0	135000.0	6750.0	135000.0	
3	0	135000.0	312682.5	29686.5	297000.0	
4	0	121500.0	513000.0	21865.5	513000.0	

5 rows × 135 columns



In [36]: `traindf.columns`

Out[36]: Index(['CNT\_CHILDREN', 'AMT\_INCOME\_TOTAL', 'AMT\_CREDIT', 'AMT\_ANNUITY', 'AMT\_GOODS\_PRICE', 'REGION\_POPULATION\_RELATIVE', 'DAYS\_BIRTH', 'DAYS\_EMPLOYED', 'DAYS\_REGISTRATION', 'DAYS\_ID\_PUBLISH', ..., 'NAME\_CONTRACT\_TYPE\_Cash loans', 'NAME\_CONTRACT\_TYPE\_Revolving loans', 'ORGANIZATION\_TYPE', 'OCCUPATION\_TYPE', 'NAME\_INCOME\_TYPE', 'WALLSMATERIAL\_MODE', 'WEEKDAY\_APPR\_PROCESS\_START', 'NAME\_TYPE\_SUITE', 'NAME\_HOUSING\_TYPE', 'NAME\_FAMILY\_STATUS'], dtype='object', length=135)

## Model Building

### 1. Xgboost (with Tuning Hyperparameter)

In [37]: `xgbpipeline = Pipeline([  
 ('scale', StandardScaler()),  
 ('clf', xgb.XGBClassifier())])`

In [39]: `# Create the parameter grid  
gbm_param_grid = {  
 'clf__learning_rate': np.arange(0.05, 1, 0.05),  
 'clf__max_depth': np.arange(3, 15, 1),  
 'clf__n_estimators': np.arange(50, 300, 50),  
 'clf__clf_colsample_bytree' : [0.2,0.4,0.6,0.8,1.0]  
}`



```
In [40]: # Perform RandomizedSearchCV
randomized_roc_auc = RandomizedSearchCV(estimator=xgbpipeline, param_distributions=gbm_param_grid,
                                         n_iter=10, scoring='roc_auc', cv=5,
                                         random_state=123, n_jobs = -2)
```

```
In [41]: # Fit the estimator
randomized_roc_auc.fit(traindf,Y_train)
```

```
Out[41]: RandomizedSearchCV(cv=5, error_score=nan,
                           estimator=Pipeline(memory=None,
                                             steps=[('scale',
                                                    StandardScaler(copy=True,
                                                                    with_mean=True,
                                                                    with_std=True)),
                                                    ('clf',
                                                     XGBClassifier(base_score=None,
                                                                    booster=None,
                                                                    colsample_bylevel
=None,
                                                                    colsample_bynode=
None,
                                                                    colsample_bytree=
None,
                                                                    gamma=None,
                                                                    gpu_id=None,
                                                                    importance_type
='gain',
                                                                    interaction_const
raints=None,
                                                                    learning_rate=
N...
                                                                    param_distributions={'clf__clf_colsample_bytrees': [0.2, 0.
4,
                                                                    0.6, 0.
8,
                                                                    1.0],
                                                                    'clf__learning_rate': array([0.05, 0.
1 , 0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 , 0.45, 0.5 , 0.55,
                                                                    0.6 , 0.65, 0.7 , 0.75, 0.8 , 0.85, 0.9 , 0.95]),
                                                                    'clf__max_depth': array([ 3,  4,  5,
6,  7,  8,  9, 10, 11, 12, 13, 14]),
                                                                    'clf__n_estimators': array([ 50, 100,
150, 200, 250])}],
                           pre_dispatch='2*n_jobs', random_state=123, refit=True,
                           return_train_score=False, scoring='roc_auc', verbose=0)
```

## Best Estimator of XGB Model

```
In [42]: # Compute metrics
print(randomized_roc_auc.best_estimator_)
```

```
Pipeline(memory=None,
          steps=[('scale',
                  StandardScaler(copy=True, with_mean=True, with_std=True)),
                 ('clf',
                  XGBClassifier(base_score=0.5, booster=None,
                                clf_colsample_bytree=0.8, colsample_bylevel=1,
                                colsample_bynode=1, colsample_bytree=1, gamma=
0,
                                gpu_id=-1, importance_type='gain',
                                interaction_constraints=None, learning_rate=0.
05,
                                max_delta_step=0, max_depth=9,
                                min_child_weight=1, missing=nan,
                                monotone_constraints=None, n_estimators=250,
                                n_jobs=0, num_parallel_tree=1,
                                objective='binary:logistic', random_state=0,
                                reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                                subsample=1, tree_method=None,
                                validate_parameters=False, verbosity=None)),
                 ],
          verbose=False)
```

```
In [43]: print(randomized_roc_auc.best_score_)
```

```
0.8447346765271921
```

```
In [44]: model_xgb_probs = randomized_roc_auc.predict_proba(traindf)
```

```
In [45]: scores = randomized_roc_auc.predict_proba(traindf)[:,-1]
```


```
fpr, tpr, thresholds = roc_curve(Y_train, scores)
```

```
roc_auc = roc_auc_score(Y_train, scores)
```

```
print("AUC of ROC Curve:", roc_auc)
```

```
AUC of ROC Curve: 0.9551377164186577
```

```
In [46]: plt.plot(fpr, tpr)
plt.title("ROC Curve XGB")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```



```
In [47]: predictions = [round(value) for value in scores]
accuracy = accuracy_score(Y_train, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```

Accuracy: 94.24%

## Predict for Test Data

```
In [85]: testdf.shape
```

Out[85]: (48744, 121)

```
In [86]: SK_ID_CURR_Col = testdf['SK_ID_CURR']
```

```
In [87]: testdf.drop('SK_ID_CURR', axis=1, inplace=True)
```

## Transform

```
In [88]: df_ohe_test = pd.get_dummies(testdf[OHE_List])
```

```
In [89]: df_le_test = testdf[LE_List].apply(lambda x: le.fit_transform(x))
```

```
In [90]: testdf.drop(categorical_columns, axis=1, inplace=True)
```

```
In [91]: testdf = pd.concat([testdf, df_ohe_test, df_le_test], axis=1)
```

```
In [92]: testdf.shape
```

```
Out[92]: (48744, 135)
```

```
In [93]: test_cols = testdf.columns.tolist()
         train_cols = traindf.columns.tolist()
```

Check if Train and Test have same set of Variables

```
In [94]: list(set(train_cols) - set(test_cols))
```

```
Out[94]: []
```

### Fit XGB to Test Data

```
In [95]: score_test = randomized_roc_auc.predict_proba(testdf)[: ,1]
```

```
In [ ]: submit_df1 = pd.DataFrame({'SK_ID_CURR':SK_ID_CURR_Col, 'TARGET': score_test})
```

```
In [ ]: submit_df1.to_csv("submit_df1.csv",index=False)
```

**The Score is 73.2**